

INFO8006 Introduction to Artificial Intelligence

Exercise session 3: Quantifying uncertainty

RECAP

Definitions

A probability is a degree of belief in a proposition.

Sample space	Ω , of possible worlds ω
Axioms	$0 \leq P(\omega) \leq 1, \quad \sum_{\omega \in \Omega} P(\omega) = 1$
Event	$A \subseteq \Omega, \quad P(A) = \sum_{\omega \in A} P(\omega)$
Random variable	$X : \Omega \rightarrow D_X$, with $X = x$ the event $\{\omega \in \Omega : X(\omega) = x\}$
Distribution	$\mathbf{P}(X)$, the numbers $P(x) = P(X = x)$ for $x \in D_X$
Joint distribution	$\mathbf{P}(X_1, \dots, X_n)$, the numbers $P(x_1, \dots, x_n)$
Conditional probability	$P(a b) = P(a, b) / P(b)$, for $P(b) > 0$
Conditional distribution	$\mathbf{P}(X Y)$, one $\mathbf{P}(X y)$ for each $y \in D_Y$

Random variables are written in upper case and their values in lower case; a boolean variable A takes the values a and $\neg a$. $P(x)$ is a number, $\mathbf{P}(X)$ a table.

Rules

Complement	$P(\neg a) = 1 - P(a)$
Union	$P(a \vee b) = P(a) + P(b) - P(a \wedge b)$
Product rule	$P(a, b) = P(a b)P(b) = P(b a)P(a)$
Sum rule, total probability	$P(x) = \sum_y P(x, y) = \sum_y P(x y)P(y)$
Chain rule	$P(x_{1:n}) = \prod_i P(x_i x_{1:i-1})$
Bayes' rule	$P(a b) = P(b a)P(a) / P(b)$
Independence, $X \perp Y$	$P(x, y) = P(x)P(y)$
Conditional independence, $X \perp Y Z$	$P(x, y z) = P(x z)P(y z)$
Naive Bayes	$P(c, e_{1:n}) = P(c) \prod_i P(e_i c)$

The rules hold for all values, and between distributions. Naive Bayes assumes the effects $e_{1:n}$ conditionally independent given the cause c . The joint distribution answers every query: select the entries consistent with the evidence, sum out the other variables, normalize.

Bayes' rule

A model is given in the *causal* direction, a *prior* $\mathbf{P}(A)$ and a *likelihood* $\mathbf{P}(B | A)$. The *diagnostic* question is answered by the *posterior* $\mathbf{P}(A | b) = \frac{1}{Z} \mathbf{P}(b | A) \mathbf{P}(A)$, with $Z = P(b) = \sum_a P(b | a)P(a)$ the probability of the *evidence*. For observations conditionally independent given A , each posterior is the prior of the next, $\mathbf{P}(A | b_{1:k+1}) = \frac{1}{Z} \mathbf{P}(b_{k+1} | A) \mathbf{P}(A | b_{1:k})$.

In session: exercises 1, 2, 3, 5.

Exercise 1 Rational beliefs (AIMA ch. 12)

Let $A, B \subseteq \Omega$ be two *events*. An agent holds the beliefs $P(A) = 0.4$ and $P(B) = 0.3$. What ranges of probabilities would it be rational for the agent to hold for the events $A \cup B$ and $A \cap B$? What if the agent also believes that $P(A \mid B) = 0.5$?

Exercise 2 Full joint distribution (AIMA ch. 12)

A monitoring program supervises the bearings of a wind turbine. A worn bearing (*fault*) tends to make the turbine vibrate (*vibration*) and to trigger the alarm of a diagnostic probe (*alarm*). The program holds the joint distribution $\mathbf{P}(V, A, F)$ below, over the boolean random variables V (vibration), A (alarm) and F (fault). Compute the following probabilities.

	v		$\neg v$	
	a	$\neg a$	a	$\neg a$
f	0.108	0.012	0.072	0.008
$\neg f$	0.016	0.064	0.144	0.576

- (a) $P(v)$
- (b) $P(f)$
- (c) $P(v \mid f)$
- (d) $P(f \mid v \vee a)$

Exercise 3 Fraud detection (AIMA ch. 12)

A fraud detection program screens card payments. It is 99 % accurate, *i.e.* it flags a fraudulent payment with probability 0.99 and a legitimate payment with probability 0.01. Fraud is rare: it strikes only 1 payment in 10 000. The program has just flagged your payment. Why is it good news, for you, that fraud is rare? What is the probability that the payment is actually fraudulent?

Exercise 4 Conditional independence (AIMA ch. 12)

For each of the following statements, either prove it is valid or give a counterexample.

- (a) If $P(a, b, c) > 0$ and $P(a \mid b, c) = P(b \mid a, c)$, then $P(a \mid c) = P(b \mid c)$.
- (b) If $P(a \mid b, c) = P(a)$, then $P(b \mid c) = P(b)$.
- (c) If $P(a \mid b) = P(a)$, then $P(a \mid b, c) = P(a \mid c)$.

Exercise 5 Bag of coins (AIMA ch. 13)

We have a bag of three biased coins a , b , and c with probabilities of coming up heads of 20 %, 60 %, and 80 %, respectively. One coin is drawn randomly from the bag (with equal probability of drawing each of the three coins), and then the coin is flipped three times to generate the outcomes X_1 , X_2 and X_3 (heads or tails).


- (a) Draw the Bayesian network corresponding to this setup and define the corresponding conditional probability tables (CPTs).
- (b) Determine which coin (a , b or c) is most likely to have been drawn from the bag if the observed tosses came out heads twice and tails once.

Exercise 6 Evolutionary model merging (AIMA ch. 13)

Evolutionary model merging builds a new language model M from two parent models A and B : each component of M is copied from one of the parents. Consider one component, and let G_x be the language it was mostly trained on in model x , with possible values e (English) or f (French). A common *hypothesis* is that the default language of a model follows from this component by a simple mechanism: the language H_x in which model x answers by default is the same as G_x with some probability s .

Furthermore, M copies the component from either parent with equal probability, and a random perturbation flips it with a small nonzero probability m .

- (a) Which of the following networks claim that $\mathbf{P}(G_M \mid G_A, G_B) = \mathbf{P}(G_M)$? Shaded variables are observed.



- (b) Which of the networks make independence claims that are consistent with the hypothesis about the default language?
- (c) Write down the CPT for the G_M node in the first network, in terms of s and m .
- (d) Suppose that $P(G_A = f) = P(G_B = f) = q$. In the first network, derive an expression for $P(G_M = f)$ in terms of s , m and q .
- (e) Over many generations of merges, we expect the distribution of the component to be the same in the parents and in the merged models. Use this to calculate the value of q . Only about 10 % of the merged models answer in French by default: explain why the hypothesis described at the beginning of this question must be wrong.

QUIZ

1. The Bayes' rule states that ...

- ☐ $P(x \mid y) = \frac{P(y \mid x)P(y)}{P(x)}$.
- ☐ $P(x \mid y) = \frac{P(y \mid x)P(x)}{P(y)}$.
- ☐ $P(y \mid x) = \frac{P(x \mid y)P(x)}{P(y)}$.
- ☐ $P(x) = \frac{P(x \mid y)P(y)}{P(x)}$.

2. Which of the following is true?

- ☐ $P(x, y \mid z) = \frac{P(x \mid y, z)}{P(y, z \mid x)}$.
- ☐ $P(x, y, z) = P(x)P(y)P(z)$.
- ☐ $P(x, y \mid z) = P(y \mid x, z)P(x \mid z)$.
- ☐ $P(x \mid y, z) = \frac{P(x, y, z)}{P(y)P(z)}$.

3. Naive Bayes model assumes ...
- ☐ conditional independence between the effects, given the cause.
 - ☐ independence between the effects.
 - ☐ independence between the effects and the cause.
 - ☐ conditional independence between each effect and the cause, given the other effects.
4. The Bayes' rule can be read as "The posterior is equal to ...
- ☐ the product of the likelihood and the prior".
 - ☐ the ratio between the posterior distribution and the evidence, times the prior".
 - ☐ the likelihood to prior ratio multiplied by the evidence".
 - ☐ the likelihood to evidence ratio multiplied by the prior".