

Introduction to Artificial Intelligence

Lecture 2: Solving problems by searching


Prof. Gilles Louppe

g.louppe@uliege.be



Today

- Planning agents
- Search problems
- Uninformed search methods
 - Depth-first search
 - Breadth-first search
 - Uniform-cost search
- Informed search methods
 - A*
 - Heuristics



Planning agents

Reflex agents

A **reflex agent** acts on the current percept alone, $a_t = \pi(e_t)$. It may keep a model of the world, but it ignores the percept history and the future consequences of its actions.

Reflexes are enough to act rationally when the environment is **fully observable**, **deterministic** and **known**.

For complex tasks, they are difficult to implement: the action with the best long-term outcome is more naturally expressed as planning.

Problem-solving agents

A **problem-solving agent** is a **goal-based agent** (Lecture 1): given goal states $G \subseteq \mathcal{S}$, it **searches** its transition model for an action sequence that reaches one of them, instead of reacting to the current percept.

Assumptions on the environment:

- **single agent**: no other agent acts on the environment;
- **fully observable**: the percept reveals the state, $e_t = s_t$;
- **deterministic**: $s_{t+1} = \text{result}(s_t, a_t)$;
- **known**: the agent knows **result**.

The agent can therefore plan the whole sequence before acting.

function SIMPLE-PROBLEM-SOLVING-AGENT(*percept*) **returns** an action

persistent: *seq*, an action sequence, initially empty

state, some description of the current world state

goal, a goal, initially null

problem, a problem formulation

state $\leftarrow$ UPDATE-STATE(*state*, *percept*)

if *seq* is empty **then**

goal $\leftarrow$ FORMULATE-GOAL(*state*)

problem $\leftarrow$ FORMULATE-PROBLEM(*state*, *goal*)

seq $\leftarrow$ SEARCH(*problem*)

if *seq* = *failure* **then return** a null action

action $\leftarrow$ FIRST(*seq*)

seq $\leftarrow$ REST(*seq*)

return *action*

Offline vs. Online solving


- **Offline**: the agent computes the whole sequence $a_{1:T}$ from s_1 , then executes it "eyes closed", ignoring its percepts $e_{1:t}$.
- **Online**: the agent recomputes what to do at each step from what it has perceived, $a_t = \pi(e_{1:t})$. That is a policy, not a plan.


Search problems

Search problems

A **search problem** is defined by

- the **states** $s \in \mathcal{S}$ of the agent and its environment, and the **initial state** $s_1 \in \mathcal{S}$;
- the **actions** $\text{actions}(s) \subseteq \mathcal{A}$ available in a state s ;
- a **transition model** $s' = \text{result}(s, a)$, the state that results from taking a in s .
 - We say that s' is a **successor** of s if some applicable action leads from s to s' .






- Together, s_1 , **actions** and **result** define the **state space**, the set of states reachable from s_1 by any sequence of actions.
 - It forms a directed graph: nodes are states, edges are actions.
 - A path is a sequence of states connected by actions.
- The **goal states** $G \subseteq \mathcal{S}$ are those in which the problem is solved. The **goal test** checks whether $s \in G$.
- The **step costs** $c(s, a, s') > 0$ give the cost of taking a in s to reach s' . The **path cost** is their sum along a path.

A **solution** is an action sequence $a_{1:T}$ such that $s_{t+1} = \text{result}(s_t, a_t)$ for $t = 1, \dots, T$ and $s_{T+1} \in G$. Its path cost is

$$\sum_{t=1}^T c(s_t, a_t, s_{t+1}).$$

An **optimal solution** is a solution of minimal path cost.

Example: Traveling in Romania




How to go from Arad to Bucharest?

- Representation of states: the city we are in.
 - $s \in \{\text{in}(\text{Arad}), \text{in}(\text{Bucharest}), \dots\}$
- Initial state = the city we start in.
 - $s_1 = \text{in}(\text{Arad})$
- Actions = Going from the current city to the cities that are directly connected to it.
 - $\text{actions}(s_1) = \{\text{go}(\text{Sibiu}), \text{go}(\text{Timisoara}), \text{go}(\text{Zerind})\}$
- Transition model = The city we arrive in after driving to it.
 - $\text{result}(\text{in}(\text{Arad}), \text{go}(\text{Zerind})) = \text{in}(\text{Zerind})$
- Goal states: being in Bucharest.
 - $G = \{\text{in}(\text{Bucharest})\}$
- Step cost: distances between cities.

Selecting a state space

The real world is absurdly **complex**.

- The **world state** includes every last detail of the environment.
- A **search state** $s \in \mathcal{S}$ keeps only the details needed for planning.




Search problems are **models**.

Example: eat-all-dots

In a maze with k dots,

- states: a position and the dots left, $s = ((x, y), d)$ with $d \in \{0, 1\}^k$;
- actions: $\mathcal{A} = \{N, S, E, W\}$;
- transition model: $\text{result}(s, a)$ moves Pacman one cell, and clears the dot it finds there;
- goal states: $G = \{((x, y), d) : d = \mathbf{0}\}$;
- step cost: 1 per move.



State space size


The maze on the right has 120 cells, 30 dots and 2 ghosts. A **world state** holds

- the position of Pacman, 120 values;
- the dots left, 2^{30} ;
- the positions of the ghosts, 12^2 ;
- the direction Pacman faces, 4;

hence $120 \times 2^{30} \times 12^2 \times 4 \approx 7 \times 10^{13}$

world states, against

$|\mathcal{S}| = 120 \times 2^{30} \approx 1.3 \times 10^{11}$ for eat-all-dots.




Tree search algorithms

We look for a path from s_1 to G in the state space.

To find one, search **grows candidate paths** from s_1 , one action at a time, until one of them reaches G .

The candidates form the **search tree**: its root is s_1 , its branches are actions, and a **node** n is a path, of last state $s(n)$ and cost $g(n)$. A state reached by several paths sits in several nodes.



function TREE-SEARCH(*problem*, *fringe*) **returns** a solution, or failure

fringe $\leftarrow$ INSERT(MAKE-NODE(INITIAL-STATE[*problem*]), *fringe*)

loop do

if *fringe* is empty **then return** failure

node $\leftarrow$ REMOVE-FRONT(*fringe*)

if GOAL-TEST(*problem*, STATE[*node*]) **then return** *node*


fringe $\leftarrow$ INSERTALL(EXPAND(*node*, *problem*), *fringe*)

Important ideas


- **Fringe** (or frontier): the leaves of the tree, generated but not expanded.
- **Expansion**: replacing a leaf by its children.
- **Strategy**: which leaf to expand next.

The tree is far too large to build. Which leaves do we expand to reach G , and how do we expand as few as possible?


(a) The initial state



(b) After expanding Arad



(c) After expanding Sibiu



Uninformed search strategies

Uninformed search strategies


Uninformed search strategies use only the information available in the problem definition. They do not know whether a state looks more promising than some other.

Strategies

- Depth-first search
- Breadth-first search
- Uniform-cost search
- Iterative deepening

Properties of search strategies


- A strategy is defined by picking the **order of expansion**.
- Strategies are evaluated along the following dimensions:
 - **Completeness**: does it always find a solution if one exists?
 - **Optimality**: does it always find the least-cost solution?
 - **Time complexity**: how long does it take to find a solution?
 - **Space complexity**: how much memory is needed to perform the search?
- Time and space complexity are measured in terms of
 - b : maximum branching factor of the search tree
 - d : depth of the least-cost solution
 - the depth of s is defined as the number of actions from the initial state to s .
 - m : maximum length of any path in the state space (may be ∞)




Depth-first search



- **Strategy:** expand the deepest node in the fringe.
- **Implementation:** fringe is a **LIFO stack**.





Properties of DFS

- **Completeness:**

- m could be infinite, so only if we prevent cycles (more on this later).

- **Optimality:**

- No, DFS finds the leftmost solution, regardless of depth or cost.


- **Time complexity:**

- May generate the whole tree (or a good part of it, regardless of d). Therefore $O(b^m)$, which might be much greater than the size of the state space!


- **Space complexity:**


- Only store siblings on path to root, therefore $O(bm)$.
- When all the descendants of a node have been visited, the node can be removed from memory.

Breadth-first search



- **Strategy:** expand the shallowest node in the fringe.
- **Implementation:** fringe is a **FIFO queue**.





Properties of BFS

- **Completeness:**

- If the shallowest goal node is at some finite depth d , BFS will eventually find it after generating all shallower nodes (provided b is finite).

- **Optimality:**

- The shallowest goal is not necessarily the optimal one.
- BFS is optimal only if the path cost is a non-decreasing function of the depth of the node.

- **Time complexity:**

- If the solution is at depth d , then the total number of nodes generated before finding this node is $b + b^2 + b^3 + \dots + b^d = O(b^d)$

- **Space complexity:**

- The number of nodes to maintain in memory is the size of the fringe, which will be the largest at the last tier. That is $O(b^d)$


(demo)

Iterative deepening


Idea: get DFS's memory with BFS's guarantees, by running DFS with depth limit 1, then 2, then 3, and so on.

It is **complete**, **optimal** when the path cost is non-decreasing with the depth, in $O(b^d)$ time and $O(bd)$ memory.

Isn't this process wastefully redundant?




Uniform-cost search



- **Strategy:** expand the cheapest node in the fringe.
- **Implementation:** fringe is a **priority queue**, using as priority the cumulative cost from s_1 to node n , that is, over the path $a_{1:k}$ that n holds,

$$g(n) = \sum_{t=1}^k c(s_t, a_t, s_{t+1}).$$



Properties of UCS

- **Completeness:**

- Yes, if the step costs are all such that $c(s, a, s') \geq \epsilon > 0$. (Why?)

- **Optimality:**

- Yes, since UCS expands nodes in order of their optimal path cost.

- **Time complexity:**

- Assume C^* is the cost of the optimal solution and that step costs are all $\geq \epsilon$.
- The "effective depth" is then roughly C^* / ϵ .
- The worst-case time complexity is $O(b^{C^* / \epsilon})$.


- **Space complexity:**

- The number of nodes to maintain is the size of the fringe, so as many as in the last tier $O(b^{C^* / \epsilon})$.

(demo)

Informed search strategies

Informed search strategies



One of the **issues of UCS** is that it explores the state space in **every direction**, without exploiting information about the (plausible) location of the goal node.

Informed search strategies aim to solve this problem by expanding nodes in the fringe in decreasing order of **desirability**.

- Greedy search
- A*


Greedy search



Heuristics

A **heuristic** (or evaluation) function $h(n)$ is:

- a function that **estimates** the cost of the cheapest path from node n to a goal state;
 - $h(n) \geq 0$ for all nodes n
 - $h(n) = 0$ for a goal state.
- is designed for a **particular** search problem.




Greedy search

- **Strategy:** expand the node n in the fringe for which $h(n)$ is the lowest.
- **Implementation:** fringe is a **priority queue**, using $h(n)$ as priority.


(a) The initial state




(b) After expanding Arad




(c) After expanding Sibiu



(d) After expanding Fagaras



$h(n)$ = straight line distance to Bucharest.



At best, greedy search takes you straight to the goal.
At worst, it is like a badly-guided BFS.

Properties of greedy search

- **Completeness:**

- No, unless we prevent cycles (more on this later).

- **Optimality:**

- No, e.g. the path via Sibiu and Fagaras is 32km longer than the path through Rimnicu Vilcea and Pitesti.

- **Time complexity:**

- $O(b^m)$, unless we have a good heuristic function.


- **Space complexity:**


- $O(b^m)$, unless we have a good heuristic function.



Shakey the Robot

- A* was first proposed in **1968** to improve robot planning.
- Goal was to navigate through a room with obstacles.






A*

- Uniform-cost orders by path cost, or **backward cost** $g(n)$
- Greedy orders by goal proximity, or **forward cost** $h(n)$
- **A*** combines the two algorithms and orders by the sum

$$f(n) = g(n) + h(n)$$

- $f(n)$ is the estimated cost of cheapest solution through n .


(a) The initial state




(b) After expanding Arad




(c) After expanding Sibiu




(d) After expanding Rimnicu Vilcea



(e) After expanding Fagaras



(f) After expanding Pitesti




Why doesn't A* stop at step (e), since Bucharest is in the fringe?

Admissible heuristics

A heuristic h is **admissible** if

$$0 \leq h(n) \leq h^*(n)$$

where $h^*(n)$ is the true cost to a nearest goal.




The Manhattan distance is admissible

Optimality of A*

Assumptions:

- A is an optimal goal node
- B is a suboptimal goal node
- h is admissible

Claim: A will exit the fringe before B .




Proof

Assume B is on the fringe, and let n be an ancestor of A on the fringe. Then

$$f(n) = g(n) + h(n) \leq g(n) + h^*(n) = g(A) = f(A) < f(B)$$


- $h(n) \leq h^*(n)$, since h is admissible;
- $g(n) + h^*(n) = g(A)$, since n is on an optimal path to A ;
- $f(A) = g(A)$ and $f(B) = g(B)$, since $h = 0$ at a goal;
- $g(A) < g(B)$, since B is suboptimal.




Therefore n expands before B . The same holds for every ancestor of A , hence A expands before B and **A^* is optimal.**

A* contours


- Assume f -costs are non-decreasing along any path.
- We can define **contour levels** t in the state space, that include all nodes n for which $f(n) \leq t$.
- A* expands every node with $f(n) < C^*$ and none with $f(n) > C^*$: it sweeps the contours outwards until it reaches the goal.



For UCS ($h(n) = 0$ for all n), bands are circular around the start.



For A* with accurate heuristics, bands stretch towards the goal.



Greedy search



UCS



A*

(demo)

Comparison of the strategies


With b the branching factor, d the depth of the least-cost solution, m the maximum path length, C^* the optimal cost and ϵ a lower bound on the step costs:

Strategy	Fringe	Complete	Optimal	Time	Space
Depth-first	LIFO stack	if no cycles	no	$O(b^m)$	$O(bm)$
Breadth-first	FIFO queue	yes	if c uniform	$O(b^d)$	$O(b^d)$
Iterative deepening	DFS, depth $1, 2, \dots$	yes	if c uniform	$O(b^d)$	$O(bd)$
Uniform-cost	priority $g(n)$	if $c \geq \epsilon$	yes	$O(b^{C^*/\epsilon})$	$O(b^{C^*/\epsilon})$
Greedy	priority $h(n)$	if no cycles	no	$O(b^m)$	$O(b^m)$
A*	priority $f(n)$	yes	if h admissible (tree)	$O(b^m)$	$O(b^m)$

Creating admissible heuristics

Most of the work in solving hard search problems optimally is in finding admissible heuristics.

Admissible heuristics can be derived from the exact solutions to **relaxed problems**, where new actions are available.



Dominance

- If h_1 and h_2 are both admissible and if $h_2(n) \geq h_1(n)$ for all n , then h_2 **dominates** h_1 and is **better** for search.
- Given any admissible heuristics h_a and h_b ,

$$h(n) = \max(h_a(n), h_b(n))$$

is also admissible and dominates h_a and h_b .

Example: the 8-puzzle

7	2	4
5		6
8	3	1

Start State

	1	2
3	4	5
6	7	8



Goal State

- $h_1(n)$, the number of misplaced tiles: $h_1 = 8$ above.
- $h_2(n)$, the sum of the Manhattan distances of the tiles to their goal positions:
 $h_2 = 3 + 1 + 2 + 2 + 2 + 3 + 3 + 2 = 18$.

Both are admissible, and h_2 **dominates** h_1 : a misplaced tile is at least one move away from its goal position, so $h_1(n) \leq h_2(n)$ everywhere. A* expands fewer nodes with h_2 .

Graph search

Graph search



A state reached by several paths sits in as many nodes: two edges per level make $2^3 = 8$ paths from A to D, for 4 states. A linear problem becomes exponential, and a cycle makes the search endless.

Graph search keeps track of the states already explored, which amounts to growing the tree directly on the state-space graph.

function GRAPH-SEARCH(*problem*, *fringe*) **returns** a solution, or failure

closed $\leftarrow$ an empty set

fringe $\leftarrow$ INSERT(MAKE-NODE(INITIAL-STATE[*problem*]), *fringe*)

loop do

if *fringe* is empty **then return** failure

node $\leftarrow$ REMOVE-FRONT(*fringe*)

if GOAL-TEST(*problem*, STATE[*node*]) **then return** *node*

if STATE[*node*] is not in *closed* **then**


 add STATE[*node*] to *closed*

fringe $\leftarrow$ INSERTALL(EXPAND(*node*, *problem*), *fringe*)

end

A* graph-search gone wrong?

- We start at *S* and *G* is a goal state.
- Which path does graph search find?



Answer

h is admissible, but it is **not consistent**: $h(A) = 4 > c(A, a, C) + h(C) = 2$.


Graph search expands S , then B ($f = 1 + 1$), then C through B ($f = 3 + 1$), which **closes** C with $g(C) = 3$. When A is expanded ($f = 1 + 4$), the cheaper path to C ($g(C) = 2$) is discarded, since C is already closed.

A* returns S, B, C, G , of cost 6, instead of S, A, C, G , of cost 5.

Consistent heuristics

A heuristic h is consistent if for every n and every successor n' generated by any action a ,

$$h(n) \leq c(n, a, n') + h(n').$$



Consequences of consistent heuristics:


- $f(n)$ is non-decreasing along any path.
- $h(n)$ is admissible.
- With a consistent heuristic, graph-search A* is optimal.

When h is not consistent

Graph search closes a state the first time it is expanded, possibly through a suboptimal path. Two ways out:

- use a **consistent** heuristic, so that a state is first expanded along its cheapest path;
- **re-expand** a closed state when it is reached with a lower f , which restores optimality for any admissible h , at the price of re-expansions.

Recap example: Super Mario



- **Task environment?**
 - the performance measure V , the environment, the actuators $\mathcal{A}$, the sensors $\mathcal{P}$?
- **Type** of environment?
 - observable, deterministic, episodic, static, discrete, single-agent, known?
- **Search problem?**
 - $s_1, \text{actions}(s), \text{result}(s, a), G, c(s, a, s')$?
- **Good heuristic** $h(n)$?



A* in action

Summary

- A **search problem** is a state space $\mathcal{S}$, an initial state s_1 , actions $\text{actions}(s)$, a transition model $\text{result}(s, a)$, goal states G and step costs c . A solution is a sequence $a_{1:T}$ reaching G , optimal when its path cost is minimal.
- A **strategy** is the order in which the fringe is expanded: LIFO for depth-first, FIFO for breadth-first, $g(n)$ for uniform-cost, $h(n)$ for greedy, $f(n) = g(n) + h(n)$ for A*.
- Only **uniform-cost** and **A*** order by cost, hence are optimal in general; breadth-first and iterative deepening return the **shallowest** goal, optimal only when the path cost never decreases with the depth, as when every step costs the same.
- h is **admissible** if $h(n) \leq h^*(n)$, which makes tree-search A* optimal, and **consistent** if $h(n) \leq c(n, a, n') + h(n')$, which makes graph-search A* optimal. A **dominating** heuristic expands fewer nodes.
- **Graph search** keeps a closed set, so a state is expanded once, along the first path reaching it.

